

Subsections

- [Initializing a Semaphore Set](#)
 - [Controlling Semaphores](#)
 - [Semaphore Operations](#)
 - [POSIX Semaphores: <semaphore.h>](#)
 - [semaphore.c: Illustration of simple semaphore passing](#)
 - [Some further example semaphore programs](#)
 - [semget.c: Illustrate the `semget\(\)` function](#)
 - [semctl.c: Illustrate the `semctl\(\)` function](#)
 - [semop\(\) Sample Program to Illustrate `semop\(\)`](#)
 - [Exercises](#)
-

IPC:Semaphores

Semaphores are a programming construct designed by E. W. Dijkstra in the late 1960s. Dijkstra's model was the operation of railroads: consider a stretch of railroad in which there is a single track over which only one train at a time is allowed. Guarding this track is a semaphore. A train must wait before entering the single track until the semaphore is in a state that permits travel. When the train enters the track, the semaphore changes state to prevent other trains from entering the track. A train that is leaving this section of track must again change the state of the semaphore to allow another train to enter. In the computer version, a semaphore appears to be a simple integer. A process (or a thread) waits for permission to proceed by waiting for the integer to become 0. The signal if it proceeds signals that this by performing incrementing the integer by 1. When it is finished, the process changes the semaphore's value by subtracting one from it.

Semaphores let processes query or alter status information. They are often used to monitor and control the availability of system resources such as shared memory segments.

Semaphores can be operated on as individual units or as elements in a set. Because System V IPC semaphores can be in a large array, they are extremely heavy weight. Much lighter weight semaphores are available in the threads library (see `man semaphore` and also Chapter [30.3](#)) and POSIX semaphores (see below briefly). Threads library semaphores must be used with mapped memory. A semaphore set consists of a control structure and an array of individual semaphores. A set of semaphores can contain up to 25 elements.

In a similar fashion to message queues, the semaphore set must be initialized using `semget()`; the semaphore creator can change its ownership or permissions using `semctl()`; and semaphore operations are performed via the `semop()` function. These are now discussed below:

Initializing a Semaphore Set

The function `semget()` initializes or gains access to a semaphore. It is prototyped by:

```
int semget(key_t key, int nsems, int semflg);
```

When the call succeeds, it returns the semaphore ID (`semid`).

The `key` argument is a access value associated with the semaphore ID.

The `nsems` argument specifies the number of elements in a semaphore array. The call fails

when `nsems` is greater than the number of elements in an existing array; when the correct count is not known, supplying 0 for this argument ensures that it will succeed.

The `semflg` argument specifies the initial access permissions and creation control flags.

The following code illustrates the `semget()` function.

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/sem.h>

...
key_t key; /* key to pass to semget() */
int semflg; /* semflg to pass to semget() */
int nsems; /* nsems to pass to semget() */
int semid; /* return value from semget() */

...

key = ...
nsems = ...
semflg = ...
if ((semid = semget(key, nsems, semflg)) == -1) {
    perror("semget: semget failed");
    exit(1); }
else
    ...
```

Controlling Semaphores

`semctl()` changes permissions and other characteristics of a semaphore set. It is prototyped as follows:

```
int semctl(int semid, int semnum, int cmd, union semun arg);
```

It must be called with a valid semaphore ID, `semid`. The `semnum` value selects a semaphore within an array by its index. The `cmd` argument is one of the following control flags:

GETVAL

-- Return the value of a single semaphore.

SETVAL

-- Set the value of a single semaphore. In this case, `arg` is taken as `arg.val`, an int.

GETPID

-- Return the `PID` of the process that performed the last operation on the semaphore or array.

GETNCNT

-- Return the number of processes waiting for the value of a semaphore to increase.

GETZCNT

-- Return the number of processes waiting for the value of a particular semaphore to reach zero.

GETALL

-- Return the values for all semaphores in a set. In this case, `arg` is taken as `arg.array`, a pointer to an array of unsigned shorts (see below).

SETALL

-- Set values for all semaphores in a set. In this case, `arg` is taken as `arg.array`, a pointer to an array of unsigned shorts.

IPC_STAT

-- Return the status information from the control structure for the semaphore set and place it in the data structure pointed to by `arg.buf`, a pointer to a buffer of type `semid_ds`.

IPC_SET

-- Set the effective user and group identification and permissions. In this case, `arg` is taken as `arg.buf`.

IPC_RMID

-- Remove the specified semaphore set.

A process must have an effective user identification of owner, creator, or superuser to perform an `IPC_SET` or `IPC_RMID` command. Read and write permission is required as for the other control commands. The following code illustrates `semctl()`.

The fourth argument `union semun arg` is optional, depending upon the operation requested. If required it is of type `union semun`, which must be *explicitly* declared by the application program as:

```

        union semun {
            int val;
            struct semid_ds *buf;
            ushort *array;
        } arg;

#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/sem.h>

union semun {
    int val;
    struct semid_ds *buf;
    ushort *array;
} arg;

int i;
int semnum = ....;
int cmd = GETALL; /* get value */

...
i = semctl(semid, semnum, cmd, arg);
if (i == -1) {
    perror("semctl: semctl failed");
    exit(1);
}
else
...

```

Semaphore Operations

`semop()` performs operations on a semaphore set. It is prototyped by:

```
int semop(int semid, struct sembuf *sops, size_t nsops);
```

The `semid` argument is the semaphore ID returned by a previous `semget()` call. The `sops` argument is a pointer to an array of structures, each containing the following information about a semaphore operation:

- The semaphore number
- The operation to be performed
- Control flags, if any.

The `sembuf` structure specifies a semaphore operation, as defined in `<sys/sem.h>`.

```
struct sembuf {
```

```

        ushort_t      sem_num;          /* semaphore number */
        short         sem_op;          /* semaphore operation */
        short         sem_flg;        /* operation flags */
};

```

The `nsops` argument specifies the length of the array, the maximum size of which is determined by the `SEMOPM` configuration option; this is the maximum number of operations allowed by a single `semop()` call, and is set to 10 by default. The operation to be performed is determined as follows:

- A positive integer increments the semaphore value by that amount.
- A negative integer decrements the semaphore value by that amount. An attempt to set a semaphore to a value less than zero fails or blocks, depending on whether `IPC_NOWAIT` is in effect.
- A value of zero means to wait for the semaphore value to reach zero.

There are two control flags that can be used with `semop()`:

IPC_NOWAIT

-- Can be set for any operations in the array. Makes the function return without changing any semaphore value if any operation for which `IPC_NOWAIT` is set cannot be performed. The function fails if it tries to decrement a semaphore more than its current value, or tests a nonzero semaphore to be equal to zero.

SEM_UNDO

-- Allows individual operations in the array to be undone when the process exits.

This function takes a pointer, `sops`, to an array of semaphore operation structures. Each structure in the array contains data about an operation to perform on a semaphore. Any process with read permission can test whether a semaphore has a zero value. To increment or decrement a semaphore requires write permission. When an operation fails, none of the semaphores is altered.

The process blocks (unless the `IPC_NOWAIT` flag is set), and remains blocked until:

- the semaphore operations can all finish, so the call succeeds,
- the process receives a signal, or
- the semaphore set is removed.

Only one process at a time can update a semaphore. Simultaneous requests by different processes are performed in an arbitrary order. When an array of operations is given by a `semop()` call, no updates are done until all operations on the array can finish successfully.

If a process with exclusive use of a semaphore terminates abnormally and fails to undo the operation or free the semaphore, the semaphore stays locked in memory in the state the process left it. To prevent this, the `SEM_UNDO` control flag makes `semop()` allocate an undo structure for each semaphore operation, which contains the operation that returns the semaphore to its previous state. If the process dies, the system applies the operations in the undo structures. This prevents an aborted process from leaving a semaphore set in an inconsistent state. If processes share access to a resource controlled by a semaphore, operations on the semaphore should not be made with `SEM_UNDO` in effect. If the process that currently has control of the resource terminates abnormally, the resource is presumed to be inconsistent. Another process must be able to recognize this to restore the resource to a consistent state. When performing a semaphore operation with `SEM_UNDO` in effect, you must also have it in effect for the call that will perform the reversing operation. When the process runs normally, the reversing operation updates the undo structure with a complementary value. This ensures that, unless the process is aborted, the values applied to the undo structure are cancel to zero. When the undo structure reaches zero, it is removed.

NOTE:Using `SEM_UNDO` inconsistently can lead to excessive resource consumption because allocated undo structures might not be freed until the system is rebooted.

The following code illustrates the `semop()` function:

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/sem.h>

...
int i;
int nsops; /* number of operations to do */
int semid; /* semid of semaphore set */
struct sembuf *sops; /* ptr to operations to perform */

...

if ((semid = semop(semid, sops, nsops)) == -1)
{
    perror("semop: semop failed");
    exit(1);
}
else
(void) fprintf(stderr, "semop: returned %d\n", i);
...
```

POSIX Semaphores: <semaphore.h>

POSIX semaphores are much lighter weight than are System V semaphores. A POSIX semaphore structure defines a single semaphore, not an array of up to twenty five semaphores. The POSIX semaphore functions are:

`sem_open()` -- Connects to, and optionally creates, a named semaphore

`sem_init()` -- Initializes a semaphore structure (internal to the calling program, so not a named semaphore).

`sem_close()` -- Ends the connection to an open semaphore.

`sem_unlink()` -- Ends the connection to an open semaphore and causes the semaphore to be removed when the last process closes it.

`sem_destroy()` -- Initializes a semaphore structure (internal to the calling program, so not a named semaphore).

`sem_getvalue()` -- Copies the value of the semaphore into the specified integer.

`sem_wait()`, `sem_trywait()` -- Blocks while the semaphore is held by other processes or returns an error if the semaphore is held by another process.

`sem_post()` -- Increments the count of the semaphore.

The basic operation of these functions is essence the same as described above, except note there are more specialised functions, here. These are not discussed further here and the reader is referred to the online `man` pages for further details.

`semaphore.c`: Illustration of simple semaphore passing

```
/* semaphore.c --- simple illustration of dijkstra's semaphore analogy
 *
 * We fork() a child process so that we have two processes running:
 * Each process communicates via a semaphore.
 * The respective process can only do its work (not much here)
 * When it notices that the semaphore track is free when it returns to 0
 * Each process must modify the semaphore accordingly
 */

#include <stdio.h>
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/sem.h>

union semun {
    int val;
    struct semid_ds *buf;
    ushort *array;
};

main()
{ int i,j;
  int pid;
  int semid; /* semid of semaphore set */
  key_t key = 1234; /* key to pass to semget() */
  int semflg = IPC_CREAT | 0666; /* semflg to pass to semget() */
  int nsems = 1; /* nsems to pass to semget() */
  int nsops; /* number of operations to do */
  struct sembuf *sops = (struct sembuf *) malloc(2*sizeof(struct sembuf));
  /* ptr to operations to perform */

  /* set up semaphore */

  (void) fprintf(stderr, "\nsemget: Setting up semaphore: semget(%#lx, %\
  %#o)\n",key, nsems, semflg);
  if ((semid = semget(key, nsems, semflg)) == -1) {
      perror("semget: semget failed");
      exit(1);
  } else
      (void) fprintf(stderr, "semget: semget succeeded: semid =\
  %d\n", semid);

  /* get child process */

  if ((pid = fork()) < 0) {
      perror("fork");
      exit(1);
  }

  if (pid == 0)
  { /* child */
      i = 0;

      while (i < 3) { /* allow for 3 semaphore sets */

          nsops = 2;

          /* wait for semaphore to reach zero */

          sops[0].sem_num = 0; /* We only use one track */
          sops[0].sem_op = 0; /* wait for semaphore flag to become zero */
          sops[0].sem_flg = SEM_UNDO; /* take off semaphore asynchronous */

          sops[1].sem_num = 0;
          sops[1].sem_op = 1; /* increment semaphore -- take control of track
          sops[1].sem_flg = SEM_UNDO | IPC_NOWAIT; /* take off semaphore */

          /* Recap the call to be made. */
```

```

(void) fprintf(stderr, "\nsemop:Child Calling semop(%d, &sops, %d) v
for (j = 0; j < nsops; j++)
{
    (void) fprintf(stderr, "\n\ttsops[%d].sem_num = %d, ", j, sops[j].
    (void) fprintf(stderr, "sem_op = %d, ", sops[j].sem_op);
    (void) fprintf(stderr, "sem_flg = %#o\n", sops[j].sem_flg);
}

/* Make the semop() call and report the results. */
if ((j = semop(semid, sops, nsops)) == -1) {
    perror("semop: semop failed");
}
else
{
    (void) fprintf(stderr, "\tsemop: semop returned %d\n", j);

    (void) fprintf(stderr, "\n\nChild Process Taking Control of
sleep(5); /* DO Nothing for 5 seconds */

    nsops = 1;

    /* wait for semaphore to reach zero */
    sops[0].sem_num = 0;
    sops[0].sem_op = -1; /* Give UP Control of track */
    sops[0].sem_flg = SEM_UNDO | IPC_NOWAIT; /* take off semaph

    if ((j = semop(semid, sops, nsops)) == -1) {
        perror("semop: semop failed");
    }
    else
        (void) fprintf(stderr, "Child Process Giving up Conti
sleep(5); /* halt process to allow parent to catch semaphor

    ++i;
}
}
else /* parent */
{ /* pid hold id of child */

    i = 0;

    while (i < 3) { /* allow for 3 semaphore sets */

        nsops = 2;

        /* wait for semaphore to reach zero */
        sops[0].sem_num = 0;
        sops[0].sem_op = 0; /* wait for semaphore flag to become zero */
        sops[0].sem_flg = SEM_UNDO; /* take off semaphore asynchronous */

        sops[1].sem_num = 0;
        sops[1].sem_op = 1; /* increment semaphore -- take control of track
        sops[1].sem_flg = SEM_UNDO | IPC_NOWAIT; /* take off semaphore */

        /* Recap the call to be made. */

        (void) fprintf(stderr, "\nsemop:Parent Calling semop(%d, &sops, %d) v
        for (j = 0; j < nsops; j++)
        {
            (void) fprintf(stderr, "\n\ttsops[%d].sem_num = %d, ", j, sops[j].
            (void) fprintf(stderr, "sem_op = %d, ", sops[j].sem_op);
            (void) fprintf(stderr, "sem_flg = %#o\n", sops[j].sem_flg);
        }

        /* Make the semop() call and report the results. */
        if ((j = semop(semid, sops, nsops)) == -1) {

```

```

        perror("semop: semop failed");
    }
    else
    {
        (void) fprintf(stderr, "semop: semop returned %d\n", j);

        (void) fprintf(stderr, "Parent Process Taking Control of Track\n");
        sleep(5); /* Do nothing for 5 seconds */

        nsops = 1;

        /* wait for semaphore to reach zero */
        sops[0].sem_num = 0;
        sops[0].sem_op = -1; /* Give UP Control of track */
        sops[0].sem_flg = SEM_UNDO | IPC_NOWAIT; /* take off semaphore */

        if ((j = semop(semid, sops, nsops)) == -1) {
            perror("semop: semop failed");
        }
        else
            (void) fprintf(stderr, "Parent Process Giving up Control of Track\n");
        sleep(5); /* halt process to allow child to catch semaphore */
    }
    ++i;
}
}
}
}

```

The key elements of this program are as follows:

- After a semaphore is created with as simple key 1234, two processes are forked.
- Each process (parent and child) essentially performs the same operations:
 - Each process accesses the same semaphore **track** (sops[0].sem_num = 0).
 - Each process waits for the **track** to become free and then attempts to take control of **track**

This is achieved by setting appropriate sops[0].sem_op values in the array.

- Once the process has control it sleeps for 5 seconds (in reality some processing would take place in place of this simple illustration)
- The process then gives up control of the **track** sops[1].sem_op = -1
- an additional sleep operation is then performed to ensure that the other process has time to access the semaphore before a subsequent (same process) semaphore read.

Note: There is no synchronisation here in this simple example and we have no control over how the OS will schedule the processes.

Some further example semaphore programs

The following suite of programs can be used to investigate interactively a variety of semaphore ideas (see exercises below).

The semaphore **must** be initialised with the semget.c program. The effects of controlling the semaphore queue and sending and receiving semaphore can be investigated with semctl.c and semop.c respectively.

semget.c: Illustrate the semget() function

```

/*
 * semget.c: Illustrate the semget() function.
 *
 * This is a simple exerciser of the semget() function. It prompts
 * for the arguments, makes the call, and reports the results.
 */

#include <stdio.h>
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/sem.h>

extern void    exit();
extern void    perror();

main()
{
    key_t  key;    /* key to pass to semget() */
    int    semflg; /* semflg to pass to semget() */
    int    nsems;  /* nsems to pass to semget() */
    int    semid;  /* return value from semget() */

    (void) fprintf(stderr,
        "All numeric input must follow C conventions:\n");
    (void) fprintf(stderr,
        "\t0x... is interpreted as hexadecimal,\n");
    (void) fprintf(stderr, "\t0... is interpreted as octal,\n");
    (void) fprintf(stderr, "\totherwise, decimal.\n");
    (void) fprintf(stderr, "IPC_PRIVATE == %#lx\n", IPC_PRIVATE);
    (void) fprintf(stderr, "Enter key: ");
    (void) scanf("%li", &key);

    (void) fprintf(stderr, "Enter nsems value: ");
    (void) scanf("%i", &nsems);
    (void) fprintf(stderr, "\nExpected flags for semflg are:\n");
    (void) fprintf(stderr, "\tIPC_EXCL = \t%#8.8o\n", IPC_EXCL);
    (void) fprintf(stderr, "\tIPC_CREAT = \t%#8.8o\n",
IPC_CREAT);
    (void) fprintf(stderr, "\towner read = \t%#8.8o\n", 0400);
    (void) fprintf(stderr, "\towner alter = \t%#8.8o\n", 0200);
    (void) fprintf(stderr, "\tgroup read = \t%#8.8o\n", 040);
    (void) fprintf(stderr, "\tgroup alter = \t%#8.8o\n", 020);
    (void) fprintf(stderr, "\tother read = \t%#8.8o\n", 04);
    (void) fprintf(stderr, "\tother alter = \t%#8.8o\n", 02);
    (void) fprintf(stderr, "Enter semflg value: ");
    (void) scanf("%i", &semflg);
    (void) fprintf(stderr, "\nsemget: Calling semget(%#lx, %
        %#o)\n", key, nsems, semflg);
    if ((semid = semget(key, nsems, semflg)) == -1) {
        perror("semget: semget failed");
        exit(1);
    } else {
        (void) fprintf(stderr, "semget: semget succeeded: semid =
%d\n",
            semid);
        exit(0);
    }
}

```

semctl.c: Illustrate the semctl() function

```

/*
 * semctl.c: Illustrate the semctl() function.
 *
 * This is a simple exerciser of the semctl() function. It lets you
 * perform one control operation on one semaphore set. It gives up

```

```

    * immediately if any control operation fails, so be careful not
    to
    * set permissions to preclude read permission; you won't be able
    to
    * reset the permissions with this code if you do.
    */

```

```

#include    <stdio.h>
#include    <sys/types.h>
#include    <sys/ipc.h>
#include    <sys/sem.h>
#include    <time.h>

```

```

struct semid_ds semid_ds;

```

```

static void do_semctl();
static void do_stat();
extern char *malloc();
extern void exit();
extern void perror();

```

```

char warning_message[] = "If you remove read permission\
    for yourself, this program will fail frequently!";

```

```

main()

```

```

{
    union semun    arg;    /* union to pass to semctl() */
    int    cmd,    /* command to give to semctl() */
        i,    /* work area */
        semid,    /* semid to pass to semctl() */
        semnum;    /* semnum to pass to semctl() */

```

```

    (void) fprintf(stderr,
        "All numeric input must follow C conventions:\n");
    (void) fprintf(stderr,
        "\t0x... is interpreted as hexadecimal,\n");
    (void) fprintf(stderr, "\t0... is interpreted as octal,\n");
    (void) fprintf(stderr, "\totherwise, decimal.\n");
    (void) fprintf(stderr, "Enter semid value: ");
    (void) scanf("%i", &semid);

```

```

    (void) fprintf(stderr, "Valid semctl cmd values are:\n");
    (void) fprintf(stderr, "\tGETALL = %d\n", GETALL);
    (void) fprintf(stderr, "\tGETNCNT = %d\n", GETNCNT);
    (void) fprintf(stderr, "\tGETPID = %d\n", GETPID);
    (void) fprintf(stderr, "\tGETVAL = %d\n", GETVAL);
    (void) fprintf(stderr, "\tGETZCNT = %d\n", GETZCNT);
    (void) fprintf(stderr, "\tIPC_RMID = %d\n", IPC_RMID);
    (void) fprintf(stderr, "\tIPC_SET = %d\n", IPC_SET);
    (void) fprintf(stderr, "\tIPC_STAT = %d\n", IPC_STAT);
    (void) fprintf(stderr, "\tSETALL = %d\n", SETALL);
    (void) fprintf(stderr, "\tSETVAL = %d\n", SETVAL);
    (void) fprintf(stderr, "\nEnter cmd: ");
    (void) scanf("%i", &cmd);

```

```

    /* Do some setup operations needed by multiple commands. */

```

```

    switch (cmd) {
        case GETVAL:
        case SETVAL:
        case GETNCNT:
        case GETZCNT:
            /* Get the semaphore number for these commands. */
            (void) fprintf(stderr, "\nEnter semnum value: ");
            (void) scanf("%i", &semnum);
            break;
        case GETALL:
        case SETALL:
            /* Allocate a buffer for the semaphore values. */
            (void) fprintf(stderr,
                "Get number of semaphores in the set.\n");
            arg.buf = &semid_ds;

```

```

do_semctl(semid, 0, IPC_STAT, arg);
if (arg.array =
    (ushort *)malloc((unsigned)
        (semid_ds.sem_nsems * sizeof(ushort)))) {
    /* Break out if you got what you needed. */
    break;
}
(void) fprintf(stderr,
    "semctl: unable to allocate space for %d values\n",
    semid_ds.sem_nsems);
exit(2);
}

/* Get the rest of the arguments needed for the specified
   command. */
switch (cmd) {
case SETVAL:
    /* Set value of one semaphore. */
    (void) fprintf(stderr, "\nEnter semaphore value: ");
    (void) scanf("%i", &arg.val);
    do_semctl(semid, semnum, SETVAL, arg);
    /* Fall through to verify the result. */
    (void) fprintf(stderr,
        "Do semctl GETVAL command to verify results.\n");
case GETVAL:
    /* Get value of one semaphore. */
    arg.val = 0;
    do_semctl(semid, semnum, GETVAL, arg);
    break;
case GETPID:
    /* Get PID of last process to successfully complete a
       semctl(SETVAL), semctl(SETALL), or semop() on the
       semaphore. */
    arg.val = 0;
    do_semctl(semid, 0, GETPID, arg);
    break;
case GETNCNT:
    /* Get number of processes waiting for semaphore value to
       increase. */
    arg.val = 0;
    do_semctl(semid, semnum, GETNCNT, arg);
    break;
case GETZCNT:
    /* Get number of processes waiting for semaphore value to
       become zero. */
    arg.val = 0;
    do_semctl(semid, semnum, GETZCNT, arg);
    break;
case SETALL:
    /* Set the values of all semaphores in the set. */
    (void) fprintf(stderr,
        "There are %d semaphores in the set.\n",
        semid_ds.sem_nsems);
    (void) fprintf(stderr, "Enter semaphore values:\n");
    for (i = 0; i < semid_ds.sem_nsems; i++) {
        (void) fprintf(stderr, "Semaphore %d: ", i);
        (void) scanf("%hi", &arg.array[i]);
    }
    do_semctl(semid, 0, SETALL, arg);
    /* Fall through to verify the results. */
    (void) fprintf(stderr,
        "Do semctl GETALL command to verify results.\n");
case GETALL:
    /* Get and print the values of all semaphores in the
       set.*/
    do_semctl(semid, 0, GETALL, arg);
    (void) fprintf(stderr,
        "The values of the %d semaphores are:\n",
        semid_ds.sem_nsems);
    for (i = 0; i < semid_ds.sem_nsems; i++)
        (void) fprintf(stderr, "%d ", arg.array[i]);

```

```

        (void) fprintf(stderr, "\n");
        break;
    case IPC_SET:
        /* Modify mode and/or ownership. */
        arg.buf = &semid_ds;
        do_semctl(semid, 0, IPC_STAT, arg);
        (void) fprintf(stderr, "Status before IPC_SET:\n");
        do_stat();
        (void) fprintf(stderr, "Enter sem_perm.uid value: ");
        (void) scanf("%hi", &semid_ds.sem_perm.uid);
        (void) fprintf(stderr, "Enter sem_perm.gid value: ");
        (void) scanf("%hi", &semid_ds.sem_perm.gid);
        (void) fprintf(stderr, "%s\n", warning_message);
        (void) fprintf(stderr, "Enter sem_perm.mode value: ");
        (void) scanf("%hi", &semid_ds.sem_perm.mode);
        do_semctl(semid, 0, IPC_SET, arg);
        /* Fall through to verify changes. */
        (void) fprintf(stderr, "Status after IPC_SET:\n");
    case IPC_STAT:
        /* Get and print current status. */
        arg.buf = &semid_ds;
        do_semctl(semid, 0, IPC_STAT, arg);
        do_stat();
        break;
    case IPC_RMID:
        /* Remove the semaphore set. */
        arg.val = 0;
        do_semctl(semid, 0, IPC_RMID, arg);
        break;
    default:
        /* Pass unknown command to semctl. */
        arg.val = 0;
        do_semctl(semid, 0, cmd, arg);
        break;
    }
    exit(0);
}

/*
 * Print indication of arguments being passed to semctl(), call
 * semctl(), and report the results. If semctl() fails, do not
 * return; this example doesn't deal with errors, it just reports
 * them.
 */
static void
do_semctl(semid, semnum, cmd, arg)
union semun arg;
int cmd,
    semid,
    semnum;
{
    register int i; /* work area */

    (void) fprintf(stderr, "\nsemctl: Calling semctl(%d, %d, %d,
",
        semid, semnum, cmd);
    switch (cmd) {
    case GETALL:
        (void) fprintf(stderr, "arg.array = %#x)\n",
            arg.array);
        break;
    case IPC_STAT:
    case IPC_SET:
        (void) fprintf(stderr, "arg.buf = %#x)\n", arg.buf);
        break;
    case SETALL:
        (void) fprintf(stderr, "arg.array = [", arg.buf);
        for (i = 0; i < semid_ds.sem_nsems; i++) {
            (void) fprintf(stderr, "%d", arg.array[i]);
            if (i < semid_ds.sem_nsems)
                (void) fprintf(stderr, ", ");
        }
    }
}

```

```

    }
    (void) fprintf(stderr, "])\n");
    break;
case SETVAL:
default:
    (void) fprintf(stderr, "arg.val = %d\n", arg.val);
    break;
}
i = semctl(semid, semnum, cmd, arg);
if (i == -1) {
    perror("semctl: semctl failed");
    exit(1);
}
(void) fprintf(stderr, "semctl: semctl returned %d\n", i);
return;
}

/*
 * Display contents of commonly used pieces of the status
structure.
 */
static void
do_stat()
{
    (void) fprintf(stderr, "sem_perm.uid = %d\n",
        semid_ds.sem_perm.uid);
    (void) fprintf(stderr, "sem_perm.gid = %d\n",
        semid_ds.sem_perm.gid);
    (void) fprintf(stderr, "sem_perm.cuid = %d\n",
        semid_ds.sem_perm.cuid);
    (void) fprintf(stderr, "sem_perm.cgid = %d\n",
        semid_ds.sem_perm.cgid);
    (void) fprintf(stderr, "sem_perm.mode = %#o, ",
        semid_ds.sem_perm.mode);
    (void) fprintf(stderr, "access permissions = %#o\n",
        semid_ds.sem_perm.mode & 0777);
    (void) fprintf(stderr, "sem_nsems = %d\n",
semid_ds.sem_nsems);
    (void) fprintf(stderr, "sem_otime = %s", semid_ds.sem_otime ?
        ctime(&semid_ds.sem_otime) : "Not Set\n");
    (void) fprintf(stderr, "sem_ctime = %s",
        ctime(&semid_ds.sem_ctime));
}

```

semop () Sample Program to Illustrate semop ()

```

/*
 * semop.c: Illustrate the semop() function.
 *
 * This is a simple exerciser of the semop() function. It lets you
 * to set up arguments for semop() and make the call. It then
reports
 * the results repeatedly on one semaphore set. You must have read
 * permission on the semaphore set or this exerciser will fail.
(It
 * needs read permission to get the number of semaphores in the set
 * and to report the values before and after calls to semop().)
 */

#include    <stdio.h>
#include    <sys/types.h>
#include    <sys/ipc.h>
#include    <sys/sem.h>

static int    ask();
extern void    exit();
extern void    free();
extern char    *malloc();
extern void    perror();

```

```

static struct semid_ds  semid_ds;                /* status of semaphore set */

static char  error_mesg1[] = "semop: Can't allocate space for %d\
    semaphore values. Giving up.\n";
static char  error_mesg2[] = "semop: Can't allocate space for %d\
    sembuf structures. Giving up.\n";

main()
{
    register int    i;    /* work area */
    int    nsops;    /* number of operations to do */
    int    semid;    /* semid of semaphore set */
    struct sembuf    *sops;    /* ptr to operations to perform */

    (void) fprintf(stderr,
        "All numeric input must follow C conventions:\n");
    (void) fprintf(stderr,
        "\t0x... is interpreted as hexadecimal,\n");
    (void) fprintf(stderr, "\t0... is interpreted as octal,\n");
    (void) fprintf(stderr, "\totherwise, decimal.\n");
    /* Loop until the invoker doesn't want to do anymore. */
    while (nsops = ask(&semid, &sops)) {
        /* Initialize the array of operations to be performed.*/
        for (i = 0; i < nsops; i++) {
            (void) fprintf(stderr,
                "\nEnter values for operation %d of %d.\n",
                i + 1, nsops);
            (void) fprintf(stderr,
                "sem_num(valid values are 0 <= sem_num < %d): ",
                semid_ds.sem_nsems);
            (void) scanf("%hi", &sops[i].sem_num);
            (void) fprintf(stderr, "sem_op: ");
            (void) scanf("%hi", &sops[i].sem_op);
            (void) fprintf(stderr,
                "Expected flags in sem_flg are:\n");
            (void) fprintf(stderr, "\tIPC_NOWAIT =\t%#6.6o\n",
                IPC_NOWAIT);
            (void) fprintf(stderr, "\tSEM_UNDO =\t%#6.6o\n",
                SEM_UNDO);
            (void) fprintf(stderr, "sem_flg: ");
            (void) scanf("%hi", &sops[i].sem_flg);
        }

        /* Recap the call to be made. */
        (void) fprintf(stderr,
            "\nsemop: Calling semop(%d, &sops, %d) with:",
            semid, nsops);
        for (i = 0; i < nsops; i++)
        {
            (void) fprintf(stderr, "\nsops[%d].sem_num = %d, ", i,
                sops[i].sem_num);
            (void) fprintf(stderr, "sem_op = %d, ", sops[i].sem_op);
            (void) fprintf(stderr, "sem_flg = %#o\n",
                sops[i].sem_flg);
        }

        /* Make the semop() call and report the results. */
        if ((i = semop(semid, sops, nsops)) == -1) {
            perror("semop: semop failed");
        } else {
            (void) fprintf(stderr, "semop: semop returned %d\n", i);
        }
    }
}

/*
 * Ask if user wants to continue.
 *
 * On the first call:
 * Get the semid to be processed and supply it to the caller.

```

```

* On each call:
* 1. Print current semaphore values.
* 2. Ask user how many operations are to be performed on the next
*    call to semop. Allocate an array of sembuf structures
*    sufficient for the job and set caller-supplied pointer to
that
*    array. (The array is reused on subsequent calls if it is big
*    enough. If it isn't, it is freed and a larger array is
*    allocated.)
*/
static
ask(semidp, sops)
int *semidp; /* pointer to semid (used only the first time) */
struct sembuf **sops;
{
    static union semun arg; /* argument to semctl */
    int i; /* work area */
    static int nsops = 0; /* size of currently allocated
        sembuf array */
    static int semid = -1; /* semid supplied by user */
    static struct sembuf *sops; /* pointer to allocated array */

    if (semid < 0) {
        /* First call; get semid from user and the current state of
            the semaphore set. */
        (void) fprintf(stderr,
            "Enter semid of the semaphore set you want to use: ");
        (void) scanf("%i", &semid);
        *semidp = semid;
        arg.buf = &semid_ds;
        if (semctl(semid, 0, IPC_STAT, arg) == -1) {
            perror("semop: semctl(IPC_STAT) failed");
            /* Note that if semctl fails, semid_ds remains filled
                with zeros, so later test for number of semaphores will
                be zero. */
            (void) fprintf(stderr,
                "Before and after values are not printed.\n");
        } else {
            if ((arg.array = (ushort *)malloc(
                (unsigned)(sizeof(ushort) * semid_ds.sem_nsems)))
                == NULL) {
                (void) fprintf(stderr, error_mesg1,
                    semid_ds.sem_nsems);
                exit(1);
            }
        }
    }
    /* Print current semaphore values. */
    if (semid_ds.sem_nsems) {
        (void) fprintf(stderr,
            "There are %d semaphores in the set.\n",
            semid_ds.sem_nsems);
        if (semctl(semid, 0, GETALL, arg) == -1) {
            perror("semop: semctl(GETALL) failed");
        } else {
            (void) fprintf(stderr, "Current semaphore values are:");
            for (i = 0; i < semid_ds.sem_nsems;
                (void) fprintf(stderr, " %d", arg.array[i++]));
            (void) fprintf(stderr, "\n");
        }
    }
    /* Find out how many operations are going to be done in the
    next
        call and allocate enough space to do it. */
    (void) fprintf(stderr,
        "How many semaphore operations do you want %s\n",
        "on the next call to semop()?");
    (void) fprintf(stderr, "Enter 0 or control-D to quit: ");
    i = 0;
    if (scanf("%i", &i) == EOF || i == 0)
        exit(0);

```

```

if (i > nsops) {
    if (nsops)
        free((char *)sops);
    nsops = i;
    if ((sops = (struct sembuf *)malloc((unsigned)(nsops *
        sizeof(struct sembuf)))) == NULL) {
        (void) fprintf(stderr, error_msg2, nsops);
        exit(2);
    }
}
*sopsp = sops;
return (i);
}

```

Exercises

Exercise 12763

Write 2 programs that will communicate **both ways** (*i.e* each process can read and write) when run concurrently via semaphores.

Exercise 12764

Modify the `semaphore.c` program to handle synchronous semaphore communication semaphores.

Exercise 12765

Write 3 programs that communicate together via semaphores according to the following specifications: `sem_server.c` -- a program that can communicate independently (on different semaphore tracks) with two clients programs. `sem_client1.c` -- a program that talks to `sem_server.c` on one track. `sem_client2.c` -- a program that talks to `sem_server.c` on another track to `sem_client1.c`.

Exercise 12766

Compile the programs `semget.c`, `semctl.c` and `semop.c` and then

- investigate and understand fully the operations of the flags (access, creation *etc.* permissions) you can set interactively in the programs.
- Use the programs to:
 - Send and receive semaphores of 3 different semaphore tracks.
 - Inquire about the state of the semaphore queue with `semctl.c`. Add/delete a few semaphores (using `semop.c` and perform the inquiry once more.
 - Use `semctl.c` to alter a semaphore on the queue.
 - Use `semctl.c` to delete a semaphore from the queue.

Dave Marshall
1/5/1999