

Subsections

- [Running UNIX Commands from C](#)
 - [execl\(\)](#)
 - [fork\(\)](#)
 - [wait\(\)](#)
 - [exit\(\)](#)
 - [Exercises](#)
-

Process Control: <stdlib.h>, <unistd.h>

A *process* is basically a single running program. It may be a ``system" program (*e.g* login, update, csh) or program initiated by the user (textedit, dbxtool or a user written one).

When UNIX runs a process it gives each process a unique number - a process ID, *pid*.

The UNIX command `ps` will list all current processes running on your machine and will list the *pid*.

The C function `int getpid()` will return the *pid* of process that called this function.

A program usually runs as a single process. However later we will see how we can make programs run as several separate communicating processes.

Running UNIX Commands from C

We can run commands from a C program just as if they were from the UNIX command line by using the `system()` function. **NOTE:** this can save us a lot of time and hassle as we can run other (proven) programs, scripts *etc.* to do set tasks.

`int system(char *string)` -- where string can be the name of a unix utility, an executable shell script or a user program. System returns the exit status of the shell. System is prototyped in <stdlib.h>

Example: Call `ls` from a program

```
main()
{ printf(`Files in Directory are:\n');
    system(`ls -l');
}
```

`system` is a call that is made up of 3 other system calls: `execl()`, `wait()` and `fork()` (which are prototyped in <unistd>)

execl()

`execl` has 5 other related functions -- see `man` pages.

`execl` stands for *execute* and *leave* which means that a process will get executed and then terminated by `execl`.

It is defined by:

```
execl(char *path, char *arg0, ..., char *argn, 0);
```

The last parameter must always be 0. It is a **NULL terminator**. Since the argument list is variable we must have some way of telling C when it is to end. The NULL terminator does this job.

where `path` points to the name of a file holding a command that is to be executed, `arg0` points to a string that is the same as `path` (or at least its last component).

`arg1 ... argn` are pointers to arguments for the command and 0 simply marks the end of the (variable) list of arguments.

So our above example could look like this also:

```
main()
{ printf("`Files in Directory are:\n");
    execl(`/bin/ls'', `ls'', ``-l'', 0);
}
```

fork()

`int fork()` turns a single process into 2 identical processes, known as the **parent** and the **child**. On success, `fork()` returns 0 to the child process and returns the process ID of the child process to the parent process. On failure, `fork()` returns -1 to the parent process, sets `errno` to indicate the error, and no child process is created.

NOTE: The child process will have its own unique PID.

The following program illustrates a simple use of `fork`, where two copies are made and run together (multitasking)

```
main()
{ int return_value;

    printf("`Forking process\n");
    fork();
    printf("`The process id is %d
        and return value is %d\n",
        getpid(), return_value);
    execl(`/bin/ls/','', `ls'', ``-l'', 0);
    printf("`This line is not printed\n");
}
```

The Output of this would be:

```
Forking process
The process id is 6753 and return value is 0
The process id is 6754 and return value is 0
two lists of files in current directory
```

NOTE: The processes have unique ID's which will be different at each run.

It also impossible to tell in advance which process will get to CPU's time --

so one run may differ from the next.

When we spawn 2 processes we can easily detect (in each process) whether it is the child or parent since fork returns 0 to the child. We can trap any errors if fork returns a -1. **i.e.:**

```
int pid; /* process identifier */

pid = fork();
if ( pid < 0 )
    { printf("`Cannot fork!!\n'");
      exit(1);
    }
if ( pid == 0 )
    { /* Child process */ ..... }
else
    { /* Parent process pid is child's pid */
      .... }

```

wait()

int wait (int *status_location) -- will force a parent process to wait for a child process to stop or terminate. wait() return the pid of the child or -1 for an error. The exit status of the child is returned to status_location.

exit()

void exit(int status) -- terminates the process which calls this function and returns the exit status value. Both UNIX and C (forked) programs can read the status value.

By convention, a status of 0 means **normal termination** any other value indicates an error or unusual occurrence. Many standard library calls have errors defined in the `sys/stat.h` header file. We can easily derive our own conventions.

A complete example of forking program is originally titled `fork.c`:

```
/* fork.c - example of a fork in a program */
/* The program asks for UNIX commands to be typed and inputted to a string*/
/* The string is then "parsed" by locating blanks etc. */
/* Each command and corresponding arguments are put in a args array */
/* execvp is called to execute these commands in child process */
/* spawned by fork() */

/* cc -o fork fork.c */

#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>

main()
{
    char buf[1024];
    char *args[64];

    for (;;) {
        /*
         * Prompt for and read a command.
         */
        printf("Command: ");

        if (gets(buf) == NULL) {
            printf("\n");
        }
    }
}

```

```

        exit(0);
    }

    /*
     * Split the string into arguments.
     */
    parse(buf, args);

    /*
     * Execute the command.
     */
    execute(args);
}

/*
 * parse--split the command in buf into
 *         individual arguments.
 */
parse(buf, args)
char *buf;
char **args;
{
    while (*buf != NULL) {
        /*
         * Strip whitespace.  Use nulls, so
         * that the previous argument is terminated
         * automatically.
         */
        while ((*buf == ' ') || (*buf == '\t'))
            *buf++ = NULL;

        /*
         * Save the argument.
         */
        *args++ = buf;

        /*
         * Skip over the argument.
         */
        while ((*buf != NULL) && (*buf != ' ') && (*buf != '\t'))
            buf++;
    }

    *args = NULL;
}

/*
 * execute--spawn a child process and execute
 *         the program.
 */
execute(args)
char **args;
{
    int pid, status;

    /*
     * Get a child process.
     */
    if ((pid = fork()) < 0) {
        perror("fork");
        exit(1);

        /* NOTE: perror() produces a short error message on the standard
         * error describing the last error encountered during a call to
         * a system or library function.
         */
    }

    /*
     * The child executes the code inside the if.
     */
    if (pid == 0) {
        execvp(*args, args);
    }
}

```

```
    perror(*args);
    exit(1);

/* NOTE: The execv() vnd execvp versions of execl() are useful when the
   number of arguments is unknown in advance;
   The arguments to execv() and execvp() are the name
   of the file to be executed and a vector of strings contain-
   ing the arguments. The last argument string must be fol-
   lowed by a 0 pointer.

   execlp() and execvp() are called with the same arguments as
   execl() and execv(), but duplicate the shell's actions in
   searching for an executable file in a list of directories.
   The directory list is obtained from the environment.
*/
}

/*
 * The parent executes the wait.
 */
while (wait(&status) != pid)
    /* empty */ ;
}
```

Exercises

Exercise 12727

Use `popen()` to pipe the `rwho` (UNIX command) output into `more` (UNIX command) in a C program.

Dave Marshall
1/5/1999